

Distributed Systems Algorithms

Distributed system must solve 6 fundamental challenges:

1. **Consensus:** Reaching agreement on a single value despite node failures or malicious actors.
 2. **Replication & Consistency:** Duplicating data for reliability while managing how updates spread.
 3. **Leader Election:** Appointing a coordinator to manage shared resources and tasks.
 4. **Fault Detection:** Monitoring node health to trigger automatic self-healing.
 5. **Clock & Ordering:** Synchronizing time or event sequences without a central clock.
-

1. Consensus and Agreement

Ensures all nodes agree on a decision (e.g. the state of a variable), even with network delays or crashes. This is essential for maintaining a consistent global state.

Example: Distributed Bank Account

- Suppose three servers maintain the balance of the same bank account: **\$100**.
- Two clients simultaneously try to **withdraw \$70**.
- Without consensus:
 - Server A approves the first transaction.
 - Server B also approves the second transaction because it hasn't seen the update from A yet.
- Result: **balance becomes \$-40**, violating correctness.

Using a **consensus algorithm**, the nodes **agree on the order of transactions**, ensuring the final balance is correct: only **one \$70 withdrawal** succeeds first, the second sees \$30 and can be rejected or queued.

Types of consensus algorithms:

- **Crash Fault Tolerant (CFT):** Handles nodes going offline (e.g., Paxos, Raft).
- **Byzantine Fault Tolerant (BFT):** Handles nodes that lie or act maliciously (e.g., PBFT).

The Byzantine Context

- **Origin (1982):** Coined by **Leslie Lamport** in "*The Byzantine Generals Problem*".
- **The Allegory:** Generals must decide to **Attack** or **Retreat**. Traitors may send conflicting messages to different generals to prevent a unanimous decision.
- **The Rule:** To tolerate f traitors, a system requires at least $3f + 1$ total nodes.
- **First Practical Implementation: PBFT (Practical Byzantine Fault Tolerance; Castro, Liskov; 1999)** demonstrated a working solution for distributed systems.
- **Modern Impact:** With **Bitcoin (2008)**, BFT became the cornerstone of blockchain technology.

Real software

- etcd
 - Apache Kafka
 - Apache ZooKeeper
 - Consul
 - RethinkDB
 - CockroachDB
 - TiDB
-

2. Replication and Consistency

Replication

Distribute data across multiple nodes while enabling **scalability** and **fault tolerance**.

- **Consistent Hashing:** Maps keys to a **ring of nodes**, so each node is responsible for a portion of the key space.
- **Replication Factor:** Each key is stored on multiple nodes (e.g., 3 nodes) to ensure **redundancy**.

- **Node Addition/Removal:** Only a fraction of keys need to be moved, minimizing disruption compared to naive hashing.

Real Software

- Cassandra
- Amazon DynamoDB
- Redis
- Riak

Consistency

- **Strong Consistency:** All replicas update immediately; users always see the latest data (High accuracy, low availability).
- **Eventual Consistency:** Replicas converge over time; the system is temporarily out of sync but eventually uniform (Low accuracy, high availability).

Key Protocols

- **Two-Phase Commit (2PC):** A strong consistency "stop-the-world" approach where all nodes must vote "Yes" to commit.
- **Sagas:** Manages eventual consistency in microservices via a sequence of local transactions and **compensating transactions** (undo) if a step fails.
- **CRDTs:** Mathematical data structures that allow concurrent updates to be merged without conflicts (e.g., Google Docs).

Real Software

- **Cassandra** (Eventual)
 - **Amazon DynamoDB** (Eventual)
 - **PostgreSQL** (Strong)
 - **MySQL** (Strong)
-

3. Leader Election & Coordination

Leader election ensures that a distributed system **operates as a unified whole** by appointing a single **coordinator (leader)**.

The leader is responsible for **serializing operations**, **managing shared resources**, and **orchestrating tasks across worker nodes**.

- **Safety:** Ensures that **at most one leader exists at any time**, preventing split-brain scenarios.
- **Liveness:** Guarantees that if the current leader fails, the system **eventually elects a new leader**, maintaining continuous operation.

Real Software

- Raft-based strong consensus: etcd, Kafka, Consul, Nomad
 - Paxos-like: ZooKeeper
 - Quorum-based / bully / ring: Redis Sentinel, Hazelcast
-

4. Fault Detection & Recovery

Continuously monitors the system to detect failures and initiate "self-healing."

- **Heartbeat Protocols:** Periodic "I am alive" signals.
- **Failover:** Shifting workloads to a healthy standby node.
- **Checkpointing:** Rolling back to the last "safe" snapshot of the global state.

Real Software

- Docker
 - Kubernetes
 - Apache Kafka
 - Redis Sentinel
-

5. Clock Synchronization & Ordering

In distributed systems, **no single universal clock** exists. Without proper ordering:

- **Logs** become inconsistent → hard to debug.
- **Transactions** may conflict → data inconsistency.
- **Causal relationships** are lost → cannot tell which event triggered another.

Physical Clocks

- Measure real-world time; use **NTP** to sync across nodes.
- Good for **logging** and **auditing**.
- Limitation: high-speed events may be **misordered** due to drift/network delay.

Logical Clocks

- Track **event order** rather than real time.

Lamport Timestamps

- Each node keeps a counter; increments on events and messages.
- Provides **total ordering**: $e1 \rightarrow e2 \rightarrow e3$.
- Limitation: cannot detect **concurrent events**.

Vector Clocks

- Each node keeps an **array of counters** (one per node).
- Compare vectors to detect:
 - $A < B \rightarrow$ A happened before B
 - $B < A \rightarrow$ B happened before A
 - Mixed → **concurrent events / conflicts**
- Useful for **causal relationships** and **conflict resolution**.

Real Software

- **Google Spanner** – Uses **TrueTime** (GPS + atomic clocks) for globally consistent timestamps across data centers.
- **Apache Cassandra** – Uses **lightweight transactions and Lamport-style logical clocks** to order updates.
- **Amazon DynamoDB** – Implements **Lamport/vector clocks** internally to maintain eventual consistency across replicas.
- **Google Docs** – Uses CRDTs and **vector clocks** to merge concurrent edits without losing causality.

References

- [Raft Simulator](#) | [Paxos Made Simple](#)
- [Practical Byzantine Fault Tolerance \(PBFT\)](#).
- [Conflict-Free Replicated Data Types \(CRDTs\)](#)
- [Leader Election in Distributed Systems](#)