


9 - First Istio architecture

In this lesson, you will learn how to use Istio instead of Eureka and Nginx in the [service-routing example](#). Indeed, Kubernetes CoreDNS can seamlessly replace Eureka for service discovery, while the Istio Gateway can replace the Nginx reverse proxy. In addition, Istio provides secure communication between pods (**mTLS**) and powerful observability out of the box.

service-routing system



Migrating to Kubernetes & Istio

To adapt the existing system to fit an Istio-enabled Kubernetes cluster, we need to perform the following actions:

1. **Remove Nginx Reverse Proxy:** Remove the `nginx` module from your parent build files.
2. **Remove Eureka Service Discovery Server:** Delete the `eureka-service` module completely.
3. **Remove Eureka Dependencies from the microservices**
 1. Remove the `Eureka` and `Spring Cloud` dependencies from `pom.xml`.
 2. Remove Eureka-related configurations from `application.yml`.
 3. Remove any `@LoadBalanced` annotations in the Java codebase (**Kubernetes and Istio will now handle load balancing natively**)
 4. Modify your routing URIs to match the new environment. Thanks to CoreDNS, each microservice can be directly reached at `http://<microservice-name>:<port>`.
4. **Compile the Applications:** Build your Java services using `mvn clean package -Dmaven.test.skip=true`.
5. **Containerize:** Build the Docker images and push them to your container registry.

Define the infrastructure

Now that the container images are ready, we will define the Kubernetes YAML manifests. Each microservice is deployed

using a `Deployment` and exposed internally via a `Service` .
Here is an example for the `bff-service` : ```yaml
apiVersion: v1 kind: Service metadata: name: bff-service labels: app: bff-service spec: ports: - port: 8080 targetPort: 8080 name: http selector: app: bff-service

```
apiVersion: apps/v1 kind: Deployment metadata: name: bff-service spec: replicas: 1 selector:
matchLabels: app: bff-service template: metadata: labels: app: bff-service spec: containers: - name:
bff-service image: mattev02/bff-service:latest ports: - containerPort: 8080 env: - name: SERVER_PORT
value: "8080" readinessProbe: httpGet: path: /actuator/health port: 8080 initialDelaySeconds: 10
periodSeconds: 10 livenessProbe: httpGet: path: /actuator/health port: 8080 initialDelaySeconds: 20
periodSeconds: 10
```

To route external traffic into the cluster, we use Istio Gateway API. A `Gateway` handles the entry point (replacing standard Ingress objects), and an `VirtualService` directs the traffic to the corresponding microservices based on the requested path.

```
```yaml
apiVersion: networking.istio.io/v1
kind: Gateway
metadata:
 name: app-gateway
spec:
 selector:
 istio: ingressgateway
 servers:
 - port:
 number: 80
 name: http
 protocol: HTTP
 hosts:
 - "*"

apiVersion: networking.istio.io/v1
kind: VirtualService
metadata:
 name: app-routing
spec:
 hosts:
 - "*"
 gateways:
 - app-gateway
 http:
 # 1. The BFF Route
 - match:
 - uri:
 prefix: /bff
```

```

 route:
 - destination:
 host: bff-service
 port:
 number: 8080

2. The Users Route
- match:
 - uri:
 prefix: /users
 route:
 - destination:
 host: user-service
 port:
 number: 8080

3. The Posts Route
- match:
 - uri:
 prefix: /posts
 route:
 - destination:
 host: post-service
 port:
 number: 8080

4. The Comments Route
- match:
 - uri:
 prefix: /comments
 route:
 - destination:
 host: comment-service
 port:
 number: 8080


5. Default catch-all route fallback
- route:
 - destination:
 host: bff-service
 port:
 number: 8080

```

## The new architecture

The final system architecture is represented below.

You can test the architecture on your own machine by following this guide: [service-mesh README](#)



### ? What are the advantages of this system compared to the initial one?

- **Reduced codebase:** Removes boilerplate code and eliminates the need for maintaining a standalone Eureka server.
- **Language Agnosticism:** While Eureka ties you heavily to the Java ecosystem, Istio works natively with applications written in any language.
- **Offloading Resilience to Infrastructure:** In a traditional Spring Boot app, you must configure libraries like Resilience4j and write fallback logic in Java. Istio manages retries, timeouts, and circuit breaking at the platform level.
- **Simplified Observability:** Metrics, tracing, and logging are captured automatically by the Envoy sidecars.
- **Encrypted Communications:** Istio transparently enforces mTLS between pods, securing internal traffic.
- **Zero-Trust Security:** Istio allows you to define strict `AuthorizationPolicies`. For example, you can write a rule stating: "Only bff-service is allowed to call user-service. Reject traffic from any other pod."