


Distributed Caching in Microservices



Caching stores frequently accessed data so microservices can serve it without repeating expensive operations. Benefits:

- **Boosts performance** – faster responses by avoiding repeated DB queries or heavy computations.
- **Enhances resiliency** – services can respond even if backend systems are temporarily unavailable.
- **Improves scalability** – reduces load on backend systems by offloading repeated requests.

Cache Patterns

1. Cache-Aside Pattern (Application-Controlled)


Very flexible and widely used. Suitable for read-heavy workloads. Application manages cache population.

Read Flow:

- Application checks the cache.
- On miss: reads from DB → stores value in cache → returns value.

Write Flow:

- Application writes to DB.
- Optionally invalidates/updates cache (TTL-based).



Pros: Full control, flexible, good for read-heavy workloads.

Cons: More complex, risk of stale cache, extra round-trips on miss.

2. Read-Through Pattern

Cache handles loading logic. Simplifies read operations, but cache may become stale if not invalidated.


Read Flow:

- Application queries cache.

- Cache fetches from DB on miss → stores → returns value.

Write Flow:

- Application writes to DB. Cache is not updated automatically.



Pros: Simplifies client code.

Cons: Cache may serve stale data, requires cache logic to fetch from DB.

3. Write-Through Pattern


Cache and DB always consistent. Writes are slower due to double-write.

Read Flow:

- Cache serves data if present.
- On miss: fetch from DB → store → return.

Write Flow:

- Application writes to cache → cache synchronously writes to DB



Pros: Strong consistency.

Cons: Slower writes.

4. Write-Behind / Write-Back Pattern


Writes are fast for the application. DB update is async → risk of data loss if cache fails.

Read Flow:

- Like read-through.

Write Flow:

- Application writes to cache → cache asynchronously updates DB.



Pros: Very fast writes, good for high-throughput systems.

Cons: Potential data loss, complex async handling.

Cache Invalidation

Cache invalidation is a “hard problem” due to the tradeoff between performance (large, long-lived cache) and consistency (fresh data).

Ensuring cached data remains accurate is challenging in distributed systems. Strategies:

- **Time-based:** Expire entries after a set period. Simple but may serve stale data.
- **Event-based:** Invalidate on data changes. Consistent but complex to implement.
- **Size-based:** Evict entries when cache exceeds memory limit. Memory-efficient, may remove relevant data.

Resources